

1 Methodology

Each search is independent, and has its own control stack and Tarjan stack. A search is started at an arbitrary node `startNode` that has not yet been considered by any other search. Each search proceeds much as in the standard Tarjans Algorithm, as long as it does not encounter a node that is part of another current search. A difficulty occurs if suspending a search would create a cycle of searches, each locked on the next. Clearly we need to take some action to ensure progress. We transfer the relevant nodes of those searches into a single search, and continue, thereby removing the blocking-cycle.

1.1 Suspending and Resuming Searches

Each node `n` includes a field `blocked: List[Search]`, storing the searches that have encountered this node and are blocked on it. When the node is completed, those searches can be resumed. Note that testing whether `n` is complete and updating `blocked` has to be done atomically. In addition, each suspended search has a field `waitingFor`, storing the node it is waiting on. We record which searches are blocked on which others in a map `suspended` from `Search` to `Search`, encapsulated in a `Suspended` object. The `Suspended` object has an operation `suspend(s: Search, n: Node)` to record that `s` is blocked on `n`.

1.2 Scheduling

The implementation uses a number of worker threads which execute searches. We use a `Scheduler` object to provide searches for workers, thereby implementing a form of task-based parallelism. The scheduler keeps track of searches that have been unblocked as a result of the blocking node becoming complete. A dormant worker can resume one of these. (Note that when a search is unblocked, the update to the `Scheduler` is done after the updates to the search itself, so that it is not resumed in an inconsistent state.) The algorithm can proceed in one of two different modes: *rooted*, where the search starts at a particular node, but the state space is not known in advance; and *unrooted*, where the state space is known in advance, and new searches can start at arbitrary nodes.[?]